EAU2 Application
Steven Horn & Arthur Armani

Introduction:

The eau2 application will be used by customers to query large data frames for results in order to analyze big data. It will be able to read large (100+ GB) schema-on-read files and load the data into a data frame using distributed arrays. The data frame will be stored in several underlying nodes hidden in the network layer that use a distributed key-value store and communicate with one another to exchange chunks of data as necessary. It will employ a 3-layer architecture (network - data frame - customer-facing) described below.

Architecture:

The system will consist of three levels.

The first level will be the network component that contains many nodes divvying up the data in a distributed key value store. This network of nodes can interact with one another to exchange or retrieve data.

The second level will consist of the data frames and distributed arrays to store the values.

The third layer will be the customer facing layer. A data analyst can input a query and receive an answer based on computations that happen in the previous two layers.

Implementation:

The Sorer class will handle reading in data from a schema-on-read file and determining a DataFrame schema. Four types of data will be supported: booleans, integers, strings, and doubles. Data will be read only once. Basic operations happening within Sorer would include determining the overall size of passed in data and splitting it up into smaller dataframes, retrieving a schema from data, and handling rows with missing values by discarding them.
- DataFrame* generate_dataframe(Schema s)

The DataFrame class will have a particular schema derived from the result of running Sorer on the data. Basic operations will include returning the schema, getting the number of rows and columns, adding rows and columns, getting and setting values, map to run through each row, and printing. A DataFrame can be constructed from an array or a scalar value as well. Upon creation a DataFrame is linked as the value to a Key and Stored in a KDFMap. The DataFrame has a Key-Chunk Map (KVStore) which

contains chunks/portions of columns. The KVStore knows which chunk belongs to which node in the network, thus it is our distributed array.

- size_t nrows(), size_t ncols()
- void add_row(Row r), void add_column(Column* col)
- const char* serialize(DataFrame* df), DataFrame* get_dataframe(const char* str)
- DataFrame* from_array(Key* key, KDFMap* kv, KVStore* kc, size_t sz, Array* from), DataFrame* from_scalar(Key* key, KDFMap* kv, KVStore* kc, <type> val)
- int get_int(size_t col, size_t row), bool get_bool(...), double get_double(...), String* get_string(...)
- void set(size_t col, size_t row, int val), set(..., bool val), set(..., double val), set(..., String* val)
- void finalize_all()
- void map(Rower r), void print()

The primary KV Store is a <Key, String> mapping where keys have a string value (name) and an index indicating which node the value belongs to, and the String is a serialized object, either a DataFrame or a Chunk. The KVStore stores chunks internally for columns so that an entire DataFrame is not stored in one node's memory, and it stores DataFrames for the application layer's use. It is also connected to a network. The network consists of various client nodes that are able to register with a server node and then exchange messages with other nodes directly. They will each be responsible for part of a distributed key-value store, so exchanging chunks of data will be necessary. These chunks of data are represented by the Chunk object and have a mapping to a home node. Data frames and messages will be able to be serialized to allow for easy travel through sockets. We can add pairs to the KV store with the put method.

- size_t index(), sockaddr_in getMyIP(), size_t port()
- void server_init(), void client_init()
- void init_sock()
- void send_m(Message* m), Message* recv_m();
- void put(Key* k, Chunk* value), put(Key* k, DataFrame* value)
- DataFrame* get(Key* k), DataFrame* getAndWait(Key* k)

The application will consist of customer-facing code that allows users to enter queries. The results of queries will be output to the user's console (or whatever front-end they are accessing the eau2 application from).

- <return type TBD> sendQuery(String query)

Use cases:

We are not quite ready for the Degrees of Linus application use case, so we will instead show off use-cases of what we have now, which is basically reading from a file to create a dataframe and mapping a key to a dataframe.

Given a sor data file data.sor:

```
Sorer s("data.sor");

DataFrame* df = s.generate_dataframe();

df->print();
```

It is that simple; the Sorer object takes the data file and parses it to infer a schema. When the generate_dataframe() method is called, the Sorer will use the schema to create a data frame and then parses the entire file to populate the data frame with rows that match the schema. You can also give the Sorer other parameters to dictate where to start reading the file and how many bytes to read in the file:

```
Sorer s("data.sor", 100, 1000);    will read the 1000 bytes of data.sor that
come after byte 100

Sorer s("data.sor", 100, -1);      will read the entire file starting at byte 100

Sorer s("data.sor", 0, -1);   will read the entire file starting at byte 0 (same
as Sorer s("data.sor"))
```

In our second milestone, we have added the distributed key-value store.

Given a key in the form of some string, we can retrieve, set, or remove a dataframe mapping:

```
DataFrame* df = kv.get(key);

SDFMap* m = new SDFMAP();

String* key1 = new String("FRAME1");

m.put(key1, new DataFrame(new Schema("")));

m.get(key1);

m.getAndWait(key1); // blocking

m.remove(key1); // removes mapping and returns DataFrame
```

In our third milestone, we have added a pseudo network running on threads. This milestone demonstrated that one node (thread) can see updates to the KV store made by another node (thread).

```
KDFMap* masterKV = new KDFMap(); // KV Store to distribute among nodes

DemoThread d1(0, 100, masterKV); // Creates a demo that produces a data frame under key "main"

DemoThread d2(1, 200, masterKV); // Creates a demo that produces a data frame under key "confirm" with the sum of the "main" data frame

DemoThread d3(2, 300, masterKV) // Creates a demo that checks whether the sum of "confirm" is equivalent to the sum of another data frame under the key "validate";
```

In our fourth milestone, we have added a real network that relies on communication over sockets. This network is capable of registering clients, sharing a directory of nodes, sending serialized messages between nodes, and thus sharing dataframes/information whenever necessary. We have not quite linked the distributed KeyValue stores to the network as this presented a significant challenge for us this week.

```
./compiledName 1 3 127.0.0.2 8084 127.0.0.1 8083 &
./compiledName 2 3 127.0.0.3 8085 127.0.0.1 8083 &
./compiledName 0 3 127.0.0.1 8083 127.0.0.1 8083
```

The above code would create a network with three nodes, the first two being clients, and the last being the server.

```
Network* n = new Network(node_info, num_nodes);

n->server_init(); OR n->client_init();
```

The above code allows for the creation of a network and the initialization of the type of node.

In the fourth milestone, the way we run the compiled executables is the same as above, but creating the network is very different. Now, the network is directly tied to the KVStore, which has also been reworked.

Now, we create a network/kvstore by doing:

```
kv = new KVStore(node_info, num_nodes, this_node, server_ip_str, server_port);

 NetworkThread n1(kv);
```

The KVStore takes in all network information and sets up the network. We then create a NetworkThread, which launches a while loop so that the node can listen for incoming requests. After the above code, you can run your application code with the created KVStore.

Status:

We needed to overhaul our implementations of both the network and of the KVStore, which took a lot of refactoring in much of the codebase. We have everything pre-milestone 3 working, including many separate unit tests. We are currently at a point where the network is almost functioning how we want it to. We are currently able to send data frames and chunks over the network, but we are running into an issue when a key-value pair isn't present yet. Because of this, the Demo example is not quite working, let alone Word Count or Linus. If you run make test, you will see the results of all unit tests and milestones 1 and 2. We have milestone 3 (non-threaded) networking attempting to run, but it fails after a little bit and gets caught in a loop. We wanted to keep it in to show how far we have gotten with the networking, but make sure that you CTRL+C to stop the loop once the feedback in the terminal stops coming in.